

Intro to R

Part 2: Functions and Objects

Prof. Bisbee

Vanderbilt University

Lecture Date: 2023/01/18

Slides Updated: 2023-01-12

Agenda

1. Recap of last lecture

- Using packages: `install.packages()` & `require()`
- Loading and manipulating data: `readRDS()` and `%>%`

2. `tidyverse` functions

- `filter` and `select`
- `summarize` and `mutate`
- `group_by`

Loading Packages & Data

- Create an `.Rmd` file and save to your `code` folder
 - Accept defaults, Save As... (with a good name), then `knit`
- Load the `tidyverse` package

```
require(tidyverse)
```

- Download `sc_debt.Rds` from [GitHub](#) and save to your `./data` folder
- Now load the data with `readRDS("[PATH TO DATA]/sc_debt.Rds")`
 - We **create** an "object" to store the data using a left-arrow: `<-`

```
df <- readRDS("../data/sc_debt.Rds")
```

- NB: `../` means "go up one folder"

Tabular Data

- Data comes in many different formats
- **Structured data**: standardized, well-defined structure, easily accessed
 - I.e., tables, databases
 - In my YouTube example, the survey we gave was **structured**
- **Unstructured data**: messy, organic, disorganized, hard to use
 - I.e., web pages, images, videos
 - In my YouTube example, the scraped HTML code of a list of recommendations was **unstructured**
- In this class, we will always be working with **structured** data...specifically "tabular data frames"
- This still requires work to prepare!

Tabular Data Frame

- AKA a "tibble"
- These are "square" (although actually rectangular)
- Rows: **units of observation** (i.e., the entities we are studying)
 - People (each row is a survey respondent, athlete, etc.)
 - Places (each row is a state, county, country, etc.)
 - Things (each row is a tweet, firm, product, etc.)
- Columns: **variables of interest** (i.e., attributes we are studying)
 - Beliefs / behaviors / etc. (i.e., where rows are people)
 - Rainfall / crimes / etc. (i.e., where rows are places)
 - Likes / profits / etc. (i.e., where rows are things)

Looking at Data

- We now have the contents of `sc_debt.Rds` stored in the object `df`
- We can look at this object directly

```
df
```

```
## # A tibble: 2,546 × 16
##   unitid instnm      stabbr grad_...1 control region preddeg
##   <int> <chr>      <chr>    <int> <chr>    <chr>    <chr>
## 1 100654 Alabama A &... AL      33375 Public  South... Bachel...
## 2 100663 University ... AL      22500 Public  South... Bachel...
## 3 100690 Amridge Uni... AL      27334 Private South... Associ...
## 4 100706 University ... AL      21607 Public  South... Bachel...
## 5 100724 Alabama Sta... AL      32000 Public  South... Bachel...
## 6 100751 The Univers... AL      23250 Public  South... Bachel...
## 7 100760 Central Ala... AL      12500 Public  South... Associ...
## 8 100812 Athens Stat... AL      19500 Public  South... Bachel...
## 9 100830 Auburn Univ... AL      24826 Public  South... Bachel...
## 10 100858 Auburn Univ... AL      21281 Public  South... Bachel...
## # ... with 2,536 more rows, 9 more variables: openadmp <int>,
## #   adm_rate <dbl>, ccbasic <int>, sat_avg <int>,
## #   md earn wne p6 <int>, ugds <int>, costt4 a <int>,
```

Looking at Data

- What is our **unit of observation**?
 - Academic institutions: each row is a single school
- What are our **variables of interest**?
 - Let's look!

```
colnames(df) # Prints the variable names
```

```
## [1] "unitid"      "instnm"      "stabbr"  
## [4] "grad_debt_mdn" "control"     "region"  
## [7] "preddeg"     "openadmp"    "adm_rate"  
## [10] "ccbasic"     "sat_avg"     "md_earn_wne_p6"  
## [13] "ugds"        "costt4_a"    "selective"  
## [16] "research_u"
```

Good Data has Codebooks!

Name	Definition
unitid	Unit ID
instnm	Institution Name
stabbr	State Abbreviation
grad_debt_mdn	Median Debt of Graduates
control	Control Public or Private
region	Census Region
preddeg	Predominant Degree Offered: Associates or Bachelors
openadmp	Open Admissions Policy: 1=Yes, 2=No, 3=No 1st time students
adm_rate	Admissions Rate: proportion of applications accepted
ccbasic	Type of institution*
sat_avg	Average SAT scores
md_earn_wne_p6	Average Earnings of Recent Graduates
ugds	Number of undergraduates
costt4_a	Average cost of attendance (tuition-grants)
selective	Institution admits fewer than 10% of applications, 1=Yes, 0=No
research_u	Institution is a research university, 1=Yes, 0=No

Manipulating the Data

- These data are cool!
- But TMI at first
- I want to know...
 1. Where is [Vanderbilt University](#)?
 2. Which school is the most selective?
 3. Which schools produce the richest grads?

Manipulating with `tidyverse`

- The code process of `tidyverse` relies on a "pipe" symbol: `%>%`
 - I don't like this name
 - I think it should be called a "chain" because it **links code together**
 - Or maybe a "do" symbol because it tells `R` what to do
 - Others refer to it as a "then" symbol, which is a little better
- The basic grammar of `R` is: object, `%>%`, verb

```
object %>% # This is the object  
function() # This is the verb
```

Manipulating with `tidyverse`

- `tidyverse` has many useful "verbs" (i.e., functions)
 - `filter()`: subsets **rows**
 - `select()`: subsets **columns**
 - `arrange()`: sorts **rows** based on **columns**
 - `summarise()`: collapses **rows**
 - `group_by()`: groups **rows** by **columns**

Manipulating: `filter()`

- So let's look at Vandy
- `filter` will select **rows** of the data based on some criteria

```
df %>%  
  filter(instnm == "Vanderbilt University") # Only select rows with  
  Vandy
```

```
## # A tibble: 1 × 16  
##   unitid instnm      stabbr grad_...1 control region preddeg  
##   <int> <chr>      <chr>    <int> <chr>   <chr>   <chr>  
## 1 221999 Vanderbilt U... TN      14962 Private South... Bachel...  
## # ... with 9 more variables: openadmp <int>, adm_rate <dbl>,  
## #   ccbasic <int>, sat_avg <int>, md_earn_wne_p6 <int>,  
## #   ugds <int>, costt4_a <int>, selective <dbl>,  
## #   research_u <dbl>, and abbreviated variable name  
## #   1grad_debt_mdn
```

Manipulating: `select()`

- Still TMI!
- I only care about the admissions rate (`adm_rate`), the SAT scores (`sat_avg`), and the future earnings (`md_earn_wne_p6`)
- `select` will select **columns**

```
df %>%  
  filter(instnm == "Vanderbilt University") %>%  
  select(instnm, adm_rate, sat_avg, md_earn_wne_p6) # Select variables  
of interest
```

```
## # A tibble: 1 × 4  
##   instnm          adm_rate sat_avg md_earn_wne_p6  
##   <chr>          <dbl>   <int>         <int>  
## 1 Vanderbilt University 0.0912    1515         53400
```

Manipulating: `arrange()`

- How does Vandy compare...?
 - to other schools in terms of SAT scores?
 - to other schools in terms of future earnings?
 - to other schools in terms of admissions rates?
- `arrange` will sort the data based on a column (ascending!)

```
df %>%  
  arrange(sat_avg) %>% # Sort data by SAT scores  
  select(instnm, sat_avg) # Only look at name and SAT scores
```

```
## # A tibble: 2,546 × 2  
##   instnm                sat_avg  
##   <chr>                <int>  
## 1 Morgan State University      737  
## 2 Saint Augustine's University  847  
## 3 Albany State University      849  
## 4 Holy Names University        851  
## 5 Livingstone College          854
```

Manipulating: `arrange()`

- Vandy is not in the bottom 10 schools

```
df %>%  
  arrange(sat_avg) %>% # Sort data by SAT scores  
  select(instnm,sat_avg) # Only look at name and SAT scores
```

```
## # A tibble: 2,546 × 2  
##   instnm                sat_avg  
##   <chr>                <int>  
## 1 Morgan State University      737  
## 2 Saint Augustine's University  847  
## 3 Albany State University      849  
## 4 Holy Names University        851  
## 5 Livingstone College          854  
## 6 Virginia Union University    855  
## 7 Manor College                861  
## 8 Saint Louis Christian College  865  
## 9 Bacone College               875  
## 10 Paine College               876  
## # ... with 2,536 more rows
```

Manipulating: `arrange()`

- Use `desc()` to order in descending values...Vandy not in top 10 either

```
df %>%  
  arrange(desc(sat_avg)) %>% # Sort data by SAT scores (descending)  
  select(instnm,sat_avg) # Only look at name and SAT scores
```

```
## # A tibble: 2,546 × 2  
##   instnm                                sat_avg  
##   <chr>                                <int>  
## 1 California Institute of Technology    1557  
## 2 Massachusetts Institute of Technology 1547  
## 3 University of Chicago                 1528  
## 4 Harvey Mudd College                   1526  
## 5 Duke University                       1522  
## 6 Franklin W Olin College of Engineering 1522  
## 7 Washington University in St Louis     1520  
## 8 Rice University                       1520  
## 9 Yale University                       1517  
## 10 Harvard University                   1517  
## # ... with 2,536 more rows
```


Manipulating: `arrange()`

- What if we look only at "selective" schools?

```
df %>%  
  filter(adm_rate < .1) %>% # Only schools who accept < 10%  
  arrange(sat_avg,adm_rate) %>% # Sort by SAT scores (ascending)  
  select(instnm,sat_avg) # Only look at name and SAT scores
```

```
## # A tibble: 25 × 2  
##   instnm                                sat_avg  
##   <chr>                                <int>  
## 1 Colby College                        1456  
## 2 Swarthmore College                  1469  
## 3 Pomona College                      1480  
## 4 Dartmouth College                  1500  
## 5 Stanford University                 1503  
## 6 Northwestern University             1506  
## 7 Columbia University in the City of New York 1511  
## 8 Brown University                   1511  
## 9 University of Pennsylvania          1511  
## 10 Vanderbilt University              1515  
## # ... with 15 more rows
```

How does Vandy compare?

- `arrange` in descending order

```
df %>%  
  filter(adm_rate < .1) %>% # Only schools who accept < 10%  
  arrange(desc(sat_avg),adm_rate) %>% # Descending SAT scores  
  select(instnm,sat_avg) # Only look at name and SAT scores
```

```
## # A tibble: 25 × 2  
##   instnm                                sat_avg  
##   <chr>                                <int>  
## 1 California Institute of Technology    1557  
## 2 Massachusetts Institute of Technology 1547  
## 3 University of Chicago                 1528  
## 4 Duke University                       1522  
## 5 Rice University                       1520  
## 6 Harvard University                   1517  
## 7 Princeton University                 1517  
## 8 Yale University                      1517  
## 9 Vanderbilt University                1515  
## 10 Columbia University in the City of New York 1511  
## # ... with 15 more rows
```

More complicated? More %>%!

- Less selective schools by SAT with debt and state

```
df %>%  
  # Less selective schools (accept 20% to 30%)  
  filter(adm_rate > .2 & adm_rate < .3) %>%  
  # Sort by state name, then by SAT scores (descending)  
  arrange(stabbr, desc(sat_avg)) %>%  
  # Only look at variables of interest  
  select(instnm, sat_avg, grad_debt_mdn, stabbr)
```

```
## # A tibble: 37 × 4  
##   instnm                sat_avg grad_...1 stabbr  
##   <chr>                <int>    <int> <chr>  
## 1 Heritage Christian University      NA      NA AL  
## 2 University of California-Santa Ba... 1370   15000 CA  
## 3 California Polytechnic State Univ... 1342   19501 CA  
## 4 University of California-Irvine    1306   15488 CA  
## 5 California Institute of the Arts      NA   27000 CA  
## 6 University of Miami                1371   17125 FL  
## 7 Georgia Institute of Technology-M... 1418   23000 GA  
## 8 Point University                   986   26000 GA  
## 9 Grinnell College                   1457   17500 IA
```

A quick aside on missingness

- Some rows have `NA` in some columns, indicating **missing data**
 - Data can be missing for many different reasons
- `NA` values will produce `NA` summaries for common functions

```
mean(c(1,2,3))
```

```
## [1] 2
```

```
mean(c(1,2,3,NA))
```

```
## [1] NA
```

- Helpers: `is.na()` and `na.rm=T`

```
mean(c(1,2,3,NA),na.rm=T)
```

```
## [1] 2
```

A quick aside on missingness

- Use `is.na()` and `filter()` to see how many schools don't report SATs

```
df %>%  
  filter(is.na(sat_avg)) %>% # Only schools that DON'T report SATs  
  select(instnm, stabbr) # Only view name and state
```

```
## # A tibble: 1,317 × 2  
##   instnm                                stabbr  
##   <chr>                                <chr>  
## 1 Amridge University                  AL  
## 2 Central Alabama Community College  AL  
## 3 Athens State University             AL  
## 4 Chattahoochee Valley Community College AL  
## 5 Coastal Alabama Community College  AL  
## 6 Gadsden State Community College     AL  
## 7 George C Wallace State Community College-Selma AL  
## 8 Heritage Christian University        AL  
## 9 Jefferson State Community College   AL  
## 10 Lurleen B Wallace Community College AL  
## # ... with 1,307 more rows
```

Stepping back

- Thus far, lots of **data**
- Not a lot of **science**
- But remember the **Research** camp!
 1. **Observation** → **Question**
 2. **Theory** → **Hypothesis**
 3. **Data Collection / Wrangling** → **Analysis**
 4. **Results** → **Conclusion**
- We have been doing lots of **Observation**!
- Do we have any good **Research questions**?

Stepping back

- **RQ**: How might admissions and SAT scores be **related**?
 - **Theory**: selective schools have stricter criteria
 - **Hypothesis**: admissions and SAT scores should be **negatively** related
- How can we test this hypothesis?

Summarizing Data: `summarise()` + `mean()`

- We can combine base R functions with `tidyverse` functions!
 - Base R: `mean()`
 - `tidyverse`: `summarise()` (aka `summarize()`)
- Overall average SAT scores

```
df %>%  
  summarise(mean_sat = mean(sat_avg, na.rm=T)) # Average SAT scores  
for entire data
```

```
## # A tibble: 1 × 1  
##   mean_sat  
##   <dbl>  
## 1    1141.
```


Summarizing Data

- Let's unpack this

```
df %>%  
  summarise(mean_sat = mean(sat_avg, na.rm=T))
```

- Create new variable `mean_sat` that contains the `mean()` of every school's average SAT score
- `na.rm=T` means we want to ignore missing data. If not?

```
df %>%  
  summarise(mean_sat = mean(sat_avg))
```

```
## # A tibble: 1 × 1  
##   mean_sat  
##   <dbl>  
## 1      NA
```

Summarizing Data

- Recall we want see if more selective schools have higher SAT scores

```
df %>%  
  filter(adm_rate < .1) %>% # Only schools who accept < 10%  
  summarise(mean_sat_LT10 = mean(sat_avg,na.rm=T)) # Average SAT
```

```
## # A tibble: 1 × 1  
##   mean_sat_LT10  
##           <dbl>  
## 1         1510.
```

```
df %>%  
  filter(adm_rate > .1) %>% # Only schools who accept > 10%  
  summarise(mean_sat_GT20 = mean(sat_avg,na.rm=T)) # Average SAT
```

```
## # A tibble: 1 × 1  
##   mean_sat_GT20  
##           <dbl>  
## 1         1135.
```

Adding / changing variables: `mutate()`

- `mutate()` creates a new variable

```
df %>%  
  mutate(newvar = 1) %>%  
  select(instnm, newvar)
```

```
## # A tibble: 2,546 × 2  
##   instnm                                newvar  
##   <chr>                                <dbl>  
## 1 Alabama A & M University              1  
## 2 University of Alabama at Birmingham    1  
## 3 Amridge University                    1  
## 4 University of Alabama in Huntsville    1  
## 5 Alabama State University              1  
## 6 The University of Alabama              1  
## 7 Central Alabama Community College      1  
## 8 Athens State University                1  
## 9 Auburn University at Montgomery        1  
## 10 Auburn University                     1
```

Object Assignment Operator: <-

- Thus far, nothing we have done has changed `df`
- Use object assignment operator `<-` to **overwrite** an existing object

```
df <- df %>%  
  mutate(adm_rate_pct = adm_rate*100)
```

- Did it work?

```
df %>%  
  summarise(adm_rate_pct = mean(adm_rate_pct, na.rm=T),  
            adm_rate = mean(adm_rate, na.rm=T))
```

```
## # A tibble: 1 × 2  
##   adm_rate_pct adm_rate  
##       <dbl>    <dbl>  
## 1       67.9    0.679
```

Logic: `ifelse()`

- 3 inputs:
 - Logical statement
 - Value if the logic is `TRUE`
 - Value if the logic is `FALSE`
- `ifelse([LOGIC],[VALUE IF TRUE],[VALUE IF FALSE])`

Logic: `ifelse()`

- Say it out loud: "Create a new variable called `selective` that records if the school is selective or not. If the admissions rate is less than 10% (0.1), record the school as `selective = 1`. Otherwise, record the school as `selective = 0`."

```
df %>%  
  mutate(selective = ifelse([LOGIC],  
                             [VALUE IF TRUE],  
                             [VALUE IF FALSE]))
```

Logic: `ifelse()`

- Say it out loud: "Create a new variable called `selective` that records if the school is selective or not. **If the admissions rate is less than 10% (0.1)**, record the school as `selective = 1`. Otherwise, record the school as `selective = 0`."

```
df %>%  
  mutate(selective = ifelse(adm_rate < 0.1,  
                             [VALUE IF TRUE],  
                             [VALUE IF FALSE]))
```

Logic: `ifelse()`

- Say it out loud: "Create a new variable called `selective` that records if the school is selective or not. If the admissions rate is less than 10% (0.1), **record the school as `selective = 1`**. Otherwise, record the school as `selective = 0`."

```
df %>%  
  mutate(selective = ifelse(adm_rate < 0.1,  
                             1,  
                             [VALUE IF FALSE]))
```


Logic: `ifelse()`

- Say it out loud: "Create a new variable called `selective` that records if the school is selective or not. If the admissions rate is less than 10% (0.1), record the school as `selective = 1`. **Otherwise, record the school as `selective = 0`.**"

```
df %>%  
  mutate(selective = ifelse(adm_rate < 0.1,  
                             1,  
                             0))
```

Logic: `ifelse()` + `mutate()`

- Remember that if we want to keep this, we need the **assignment operator** `<-`

```
df <- df %>%  
  mutate(selective = ifelse(adm_rate < 0.1,  
                             1,  
                             0))
```

Quick Test

- Create a new variable `big` that is `1` if a school has more than 10,000 undergrads and `0` otherwise

```
# INSERT CODE HERE
```

Summarizing Data: `group_by()`

- One final `tidyverse` function: `group_by()`
- Let's use the newly created `selective` variable which is either 1 or 0

```
df %>%  
  select(instnm,selective,adm_rate)
```

```
## # A tibble: 2,546 × 3  
##   instnm                selective adm_rate  
##   <chr>                <dbl>    <dbl>  
## 1 Alabama A & M University      0     0.918  
## 2 University of Alabama at Birmingham      0     0.737  
## 3 Amridge University           NA     NA  
## 4 University of Alabama in Huntsville      0     0.826  
## 5 Alabama State University          0     0.969  
## 6 The University of Alabama          0     0.827  
## 7 Central Alabama Community College      NA     NA  
## 8 Athens State University          NA     NA  
## 9 Auburn University at Montgomery        0     0.904  
## 10 Auburn University             0     0.807  
## # ... with 2,536 more rows
```

Summarizing Data: `group_by()`

- Instead of running two separate `filter()` commands, use `group_by()`

```
df %>%  
  # Group the data by selective (either 1 or 0)  
  group_by(selective) %>%  
  # Calculate average SAT for each group  
  summarise(mean_sat = mean(sat_avg, na.rm=T))
```

```
## # A tibble: 3 × 2  
##   selective mean_sat  
##   <dbl>     <dbl>  
## 1      0     1135.  
## 2      1     1510.  
## 3     NA      NaN
```

Results

- Do more selective schools have higher SAT scores?
- Yes
- This **Result** **confirms** our **Hypothesis** and **answers** our **Research Question**

Conclusion

- What we've done today is a microcosm of data science
 1. Opened **data** (`readRDS`)
 2. Looked at **data** (`tidyverse` + `select()`, `filter()`, `arrange()`)
 3. Generated **hypotheses** (Admissions versus SAT scores)
 4. **Tested hypotheses** (`summarise()` + `mean()`)

Advanced Logic: `filter()`

If no time, jump to end

- `filter()` command with other logical operators
 - `>`, `<`: greater than, less than (`>=`, `<=`)
 - `!`: not (i.e., `!=` means "not equal to")
 - `&`: and
 - `|`: or

```
df %>%  
  # Schools EXCEPT Vandy  
  filter(instnm != "Vanderbilt University") %>%  
  select(instnm, stabbr, adm_rate, sat_avg)
```

```
## # A tibble: 2,545 × 4  
##   instnm                stabbr adm_r...1 sat_avg  
##   <chr>                <chr>    <dbl>    <int>  
## 1 Alabama A & M University AL      0.918     939  
## 2 University of Alabama at Birmingham AL      0.737    1234  
## 3 Amridge University AL      NA        NA  
## 4 University of Alabama in Huntsvil... AL      0.826    1319  
## 5 Alabama State University AL      0.969     946  
## 6 The University of Alabama AL      0.827    1261
```


Advanced Logic: `str_detect()`

- `filter()` command with other functions
 - `str_detect([VAR],[PATTERN])`: detect a string
 - `grepl([PATTERN],[VAR])`: also detects a string

```
df %>%  
  filter(str_detect(instnm, "Vanderbilt")) %>%  
  select(instnm, stabbr, adm_rate, sat_avg)
```

```
## # A tibble: 1 × 4  
##   instnm          stabbr adm_rate sat_avg  
##   <chr>          <chr>    <dbl>   <int>  
## 1 Vanderbilt University TN      0.0912   1515
```

Advanced Logic: `str_detect()`

- String detection is case sensitive!

```
df %>%  
  filter(str_detect(instnm, "VAND")) %>%  
  select(instnm, stabbr, adm_rate, sat_avg)
```

```
## # A tibble: 0 × 4  
## #   ... with 4 variables: instnm <chr>, stabbr <chr>,  
## #     adm_rate <dbl>, sat_avg <int>
```

```
df %>%  
  filter(str_detect(instnm, "anderbil")) %>%  
  select(instnm, stabbr, adm_rate, sat_avg)
```

```
## # A tibble: 1 × 4  
##   instnm          stabbr adm_rate sat_avg  
##   <chr>          <chr>    <dbl>   <int>  
## 1 Vanderbilt University TN      0.0912   1515
```

Advanced Logic: & (and), | (or)

```
df %>%  
  filter(str_detect(instnm, "Colorado")) %>%  
  select(instnm, stabbr, adm_rate, sat_avg)
```

```
## # A tibble: 12 × 4  
##   instnm                stabbr adm_r...1 sat_avg  
##   <chr>                <chr>    <dbl>    <int>  
## 1 University of Colorado Denver/Ans... CO      0.673    1124  
## 2 University of Colorado Colorado S... CO      0.872    1136  
## 3 University of Colorado Boulder      CO      0.784    1276  
## 4 Colorado Christian University      CO      NA        NA  
## 5 Colorado College                  CO      0.135     NA  
## 6 Colorado School of Mines          CO      0.531    1342  
## 7 Colorado State University-Fort Co... CO      0.814    1204  
## 8 Colorado Mesa University          CO      0.782    1063  
## 9 University of Northern Colorado    CO      0.908    1096  
## 10 Colorado State University Pueblo   CO      0.930    1047  
## 11 Western Colorado University        CO      0.842    1114  
## 12 Colorado State University-Global ... CO      0.986    1048  
## # ... with abbreviated variable name 1adm_rate
```

Advanced Logic: & (and), | (or)

```
df %>%  
  filter(grepl("Colorado",instnm) & grepl(' of ',instnm)) %>%  
  select(instnm,stabbr,adm_rate,sat_avg)
```

```
## # A tibble: 5 × 4  
##   instnm                stabbr adm_r...1 sat_avg  
##   <chr>                <chr>    <dbl>    <int>  
## 1 University of Colorado Denver/Ansc... CO      0.673    1124  
## 2 University of Colorado Colorado Sp... CO      0.872    1136  
## 3 University of Colorado Boulder        CO      0.784    1276  
## 4 Colorado School of Mines              CO      0.531    1342  
## 5 University of Northern Colorado        CO      0.908    1096  
## # ... with abbreviated variable name 1adm_rate
```

Advanced Logic: & (and), | (or)

```
df %>%  
  filter(grepl("Colorado",instnm) | grepl('Vermont',instnm)) %>%  
  select(instnm,stabbr,adm_rate,sat_avg)
```

```
## # A tibble: 16 × 4  
##   instnm                                stabbr adm_r...1 sat_avg  
##   <chr>                                <chr>    <dbl>    <int>  
## 1 University of Colorado Denver/Ans... CO      0.673    1124  
## 2 University of Colorado Colorado S... CO      0.872    1136  
## 3 University of Colorado Boulder      CO      0.784    1276  
## 4 Colorado Christian University       CO      NA        NA  
## 5 Colorado College                   CO      0.135     NA  
## 6 Colorado School of Mines            CO      0.531    1342  
## 7 Colorado State University-Fort Co... CO      0.814    1204  
## 8 Colorado Mesa University            CO      0.782    1063  
## 9 University of Northern Colorado     CO      0.908    1096  
## 10 Colorado State University Pueblo   CO      0.930    1047  
## 11 Western Colorado University        CO      0.842    1114  
## 12 Community College of Vermont       VT      NA        NA  
## 13 Northern Vermont University        VT      0.778     NA  
## 14 Vermont Technical College          VT      0.670     NA  
## 15 University of Vermont              VT      0.673    1287
```

Advanced Logic: & (and), | (or)

```
df %>%  
  filter((grepl("Colorado",instnm) | grepl('Vermont',instnm)) &  
grepl(' of ',instnm)) %>%  
  select(instnm,stabbr,adm_rate,sat_avg)
```

```
## # A tibble: 7 × 4  
##   instnm          stabbr adm_r...1 sat_avg  
##   <chr>          <chr>    <dbl>   <int>  
## 1 University of Colorado Denver/Ansc... CO      0.673    1124  
## 2 University of Colorado Colorado Sp... CO      0.872    1136  
## 3 University of Colorado Boulder        CO      0.784    1276  
## 4 Colorado School of Mines              CO      0.531    1342  
## 5 University of Northern Colorado        CO      0.908    1096  
## 6 Community College of Vermont          VT      NA        NA  
## 7 University of Vermont                  VT      0.673    1287  
## # ... with abbreviated variable name 1adm_rate
```

Advanced Logic: & (and), | (or)

- & can be separated into multiple `filter()` commands

```
df %>%  
  filter((grepl("Colorado",instnm) | grepl('Vermont',instnm))) %>%  
  filter(grepl(' of ',instnm)) %>%  
  select(instnm,stabbr,adm_rate,sat_avg)
```

```
## # A tibble: 7 × 4  
##   instnm                stabbr adm_r...1 sat_avg  
##   <chr>                <chr>    <dbl>    <int>  
## 1 University of Colorado Denver/Ansc... CO      0.673    1124  
## 2 University of Colorado Colorado Sp... CO      0.872    1136  
## 3 University of Colorado Boulder        CO      0.784    1276  
## 4 Colorado School of Mines              CO      0.531    1342  
## 5 University of Northern Colorado        CO      0.908    1096  
## 6 Community College of Vermont          VT      NA        NA  
## 7 University of Vermont                 VT      0.673    1287  
## # ... with abbreviated variable name 1adm_rate
```

Advanced Logic: & (and), | (or)

- | can be moved into the `str_detect()` or `grepl()` commands

```
df %>%  
  filter(grepl("Colorado|Vermont", instnm)) %>%  
  filter(grepl(' of ', instnm)) %>%  
  select(instnm, stabbr, adm_rate, sat_avg)
```

```
## # A tibble: 7 × 4  
##   instnm                stabbr adm_r...1 sat_avg  
##   <chr>                <chr>    <dbl>    <int>  
## 1 University of Colorado Denver/Ansc... CO      0.673    1124  
## 2 University of Colorado Colorado Sp... CO      0.872    1136  
## 3 University of Colorado Boulder        CO      0.784    1276  
## 4 Colorado School of Mines              CO      0.531    1342  
## 5 University of Northern Colorado       CO      0.908    1096  
## 6 Community College of Vermont          VT      NA        NA  
## 7 University of Vermont                 VT      0.673    1287  
## # ... with abbreviated variable name 1adm_rate
```


Quick Test

- Filter schools from Texas with the word "community" in their name

```
# INSERT CODE HERE
```

Advanced Logic: `select()`

- `select` can be paired with `matches()` or `contains()` for similar flexibility (equivalent to `str_detect()` or `grep1()` for `filter()`)

```
df %>%  
  select(contains('inst'))
```

```
## # A tibble: 2,546 × 1  
##   instnm  
##   <chr>  
## 1 Alabama A & M University  
## 2 University of Alabama at Birmingham  
## 3 Amridge University  
## 4 University of Alabama in Huntsville  
## 5 Alabama State University  
## 6 The University of Alabama  
## 7 Central Alabama Community College  
## 8 Athens State University  
## 9 Auburn University at Montgomery  
## 10 Auburn University  
## # ... with 2,536 more rows
```

Advanced Logic: `select()`

- `matches` can work with `|`

```
df %>%  
  select(!matches('_|inst'))
```

```
## # A tibble: 2,546 × 9  
##   unitid stabbr control region    preddeg opena...1 ccbasic  
##   <int> <chr>   <chr>   <chr>    <chr>      <int>   <int>  
## 1 100654 AL      Public Southeast Bachelor...      2      18  
## 2 100663 AL      Public Southeast Bachelor...      2      15  
## 3 100690 AL      Private Southeast Associate      1      20  
## 4 100706 AL      Public Southeast Bachelor...      2      16  
## 5 100724 AL      Public Southeast Bachelor...      2      19  
## 6 100751 AL      Public Southeast Bachelor...      2      15  
## 7 100760 AL      Public Southeast Associate      1       2  
## 8 100812 AL      Public Southeast Bachelor...     NA      22  
## 9 100830 AL      Public Southeast Bachelor...      2      18  
## 10 100858 AL      Public Southeast Bachelor...      2      15  
## # ... with 2,536 more rows, 2 more variables: ugds <int>,  
## #   selective <dbl>, and abbreviated variable name  
## #   1openadmp
```

Advanced Logic: `select()`

- `select` can also work with `where` to find classes

```
df %>%  
  select(where(is.numeric))
```

```
## # A tibble: 2,546 × 12  
##   unitid grad_deb...1 opena...2 adm_r...3 ccbasic sat_avg md_ea...4  
##   <int>      <int>      <int>      <dbl>      <int>      <int>      <int>  
## 1 100654      33375         2    0.918         18        939      25200  
## 2 100663      22500         2    0.737         15       1234      35100  
## 3 100690      27334         1    NA           20        NA      30700  
## 4 100706      21607         2    0.826         16       1319      36200  
## 5 100724      32000         2    0.969         19        946      22600  
## 6 100751      23250         2    0.827         15       1261      37400  
## 7 100760      12500         1    NA           2        NA      23100  
## 8 100812      19500        NA    NA           22        NA      33400  
## 9 100830      24826         2    0.904         18       1082      30100  
## 10 100858      21281         2    0.807         15       1300      39500  
## # ... with 2,536 more rows, 5 more variables: ugds <int>,  
## #   costt4_a <int>, selective <dbl>, research_u <dbl>,  
## #   adm_rate_pct <dbl>, and abbreviated variable names  
## #   1grad debt mdn, 2openadmp, 3adm rate, 4md earn wne p6
```

Quick Test

- Filter to only schools in California and select only character columns

```
# INSERT CODE HERE
```

Quiz & Homework

If time, jump to advanced

- Go to Brightspace and take the **3rd** quiz
 - The password to take the quiz is ####
- **Homework:**
 1. Work through Intro_to_R_Part2_hw.Rmd