

Data Wrangling in R

Homework

Prof. Bisbee

Due Date: 2023-02-01

Univariate Data Analysis

Univariate is pretty much what it sounds like: one variable. When undertaking univariate data analysis, we need first and foremost to figure what type of variable it is that we're working with. Once we do that, we can choose the appropriate use of the variable, either as an outcome or as a possible predictor.

Motivating Question

We'll be working with data from every NBA player who was active during the 2018-19 season.

Here's the data:

```
require(tidyverse)
nba<-read_rds("../data/nba_players_2018.Rds")
```

This data contains the following variables:

Codebook for NBA Data

Name	Definition
namePlayer	Player name
idPlayer	Unique player id
slugSeason	Season start and end
numberPlayerSeason	Which season for this player
isRookie	Rookie season, true or false
slugTeam	Team short name
idTeam	Unique team id
gp	Games Played
gs	Games Started
fgm	Field goals made
fga	Field goals attempted
pctFG	Percent of field goals made

Name	Definition
fg3m	3 point field goals made
fg3a	3 point field goals attempted
pctFG3	Percent of 3 point field goals made
pctFT	Free Throw percentage
fg2m	2 point field goals made
fg2a	2 point field goals attempted
pctFG2	Percent of 2 point field goals made
agePlayer	Player age
minutes	Minutes played
ftm	Free throws made
fta	Free throws attempted
oreb	Offensive rebounds
dreb	Defensive rebounds
treb	Total rebounds
ast	Assists
blk	Blocks
tov	Turnovers
pf	Personal fouls
pts	Total points
urlNBAAPI	Source url

We're interested in the following questions:

- Do certain colleges produce players that have more field goals? What about free throw percentage above a certain level? Are certain colleges in the east or the west more likely to produce higher scorers? How does this vary as a player has more seasons?

To answer these questions we need to look at the following variables:

- Field goals
- Free throw percentage above .25
- Colleges
- Player seasons
- Region

We're going to go through a pretty standard set of steps for each variable. First, examine some cases. Second, based on our examination, we'll try either a plot or a table. Once we've seen the plot or the table, we'll think a bit about ordering, and then choose an appropriate measure of central tendency, and maybe variation.

Types of Variables

It's really important to understand the types of variables you're working with. Many times analysts are indifferent to this step particularly with larger datasets. This can lead to a great deal of confusion down the road. Below are the variable types we'll be working with this semester and the definition of each.

Continuous Variables

A continuous variable can theoretically be subdivided at any arbitrarily small measure and can still be identified. You may have encountered further subdivision of continuous variables into "interval" or "ratio" data in other classes. We RARELY use these distinctions in practice. The distinction between a continuous and a categorical variable is hugely consequential, but the distinction between interval and ratio is not really all that important in practice.

The mean is the most widely used measure of central tendency for a continuous variable. If the distribution of the variable isn't very symmetric or there are large outliers, then the median is a much better measure of central tendency.

Categorical Variables

A categorical variable divides the sample up into a set of mutually exclusive and exhaustive categories. Mutually exclusive means that each case can only be one, and exhaustive means that the categories cover every possible option. Categorical is sort of the "top" level classification for variables of this type. Within the broad classification of categorical there are multiple types of other variables.

Categorical: ordered

an ordered categorical variable has— you guessed it— some kind of sensible order that can be applied. For instance, the educational attainment of an individual: high school diploma, associates degree, bachelor's degree, graduate degree— is an ordered categorical variable.

Ordered categorical variables should be arranged in the order of the variable, with proportions or percentages associated with each order. The mode, or the category with the highest proportion, is a reasonable measure of central tendency, but with fewer than ten categories the analyst should generally just show the proportion in each category.

Categorical: ordered, binary

An ordered binary variable has just two levels, but can be ordered. For instance, is a bird undertaking its first migration: yes or no? A "no" means that the bird has more than one.

The mean of a binary variable is exactly the same thing as the proportion of the sample with that characteristic. So, the mean of a binary variable for "first migration" where 1="yes" will give the proportion of birds migrating for the first time.

An ordered binary variable coded as 0 or 1 can be summarized using the mean which is the same thing as the proportion of the sample with that characteristic.

Categorical: unordered

An unordered categorical variable has no sensible ordering that can be applied. Think about something like college major. There's no "number" we might apply to philosophy that has any meaningful distance from a number we might apply to chemical engineering.

Unlike an ordered variable, an unordered categorical variable should be ordered in terms of the proportions falling into each of the categories. As with an unordered variable, it's best just to show the proportions in each category for variables with less than ten levels. The mode is a reasonable single variable summary of an unordered categorical variable.

Categorical: unordered, binary

This kind of variable has no particular order, but can be just binary. A “1” means that the case has that characteristic, a “0” means the case does not have that characteristic. For instance, whether a tree is deciduous or not.

An unordered binary variable coded as 0 or 1 can also be summarized by the mean, which is the same thing as the proportion of the sample with that characteristic.

Formats for categorical variables

In R, categorical variables CAN be stored as text, numbers or even logicals. Don't count on the data to help you out— you as the analyst need to figure this out.

Factors

We probably need to talk about factors. In R, a factor is a way of storing categorical variables. The factor provides additional information, including an ordering of the variable and a number assigned to each “level” of the factor. A categorical variable is a general term that's understood across statistics. A factor variable is a specific R term. Most of the time it's best not to have a categorical variable structured as a factor unless you know you want it to be a factor. More on this later . . .

The Process: #TrustTheProcess

I'm going to walk you through how an analyst might typically decide what type of variables they're working with. It generally works like this:

1. Take a look at a few observations and form a guess as to what type of variable it is.
2. Based on that guess, create an appropriate plot or table.
3. If the plot or table looks as expected, calculate some summary measures. If not, go back to 1.

“Glimpse” to start: what's in here anyway?

The first thing we're going to do with any dataset is just to take a quick look. We can call the data itself, but that will just show the first few cases and the first few variables. Far better is the `glimpse` command, which shows us all variables and the first few observations for all of the variables. Here's a link to the codebook for this dataset:

The six variables we're going to think about are field goals, free throw percentage, seasons played, rookie season, college attended, and conference played in.

```
glimpse(nba)
```

```
## Rows: 530
## Columns: 37
## $ namePlayer      <chr> "LaMarcus Aldridge", "Quincy Acy", "Steven Adams", ...
## $ idPlayer        <dbl> 200746, 203112, 203500, 203518, 1628389, 1628959, 1...
## $ slugSeason      <chr> "2018-19", "2018-19", "2018-19", "2018-19", "2018-1...
## $ numberPlayerSeason <dbl> 12, 6, 5, 2, 1, 0, 0, 0, 0, 0, 8, 5, 4, 3, 1, 1, 1,...
## $ isRookie        <lgl> FALSE, FALSE, FALSE, FALSE, FALSE, TRUE, TRUE, TRUE...
## $ slugTeam        <chr> "SAS", "PHX", "OKC", "OKC", "MIA", "CHI", "UTA", "C...
## $ idTeam          <dbl> 1610612759, 1610612756, 1610612760, 1610612760, 161...
## $ gp              <dbl> 81, 10, 80, 31, 82, 10, 38, 19, 34, 7, 81, 72, 43, ...
## $ gs              <dbl> 81, 0, 80, 2, 28, 1, 2, 3, 1, 0, 81, 72, 40, 4, 80,...
## $ fgm             <dbl> 684, 4, 481, 56, 280, 13, 67, 11, 38, 3, 257, 721, ...
## $ fga             <dbl> 1319, 18, 809, 157, 486, 39, 178, 36, 110, 10, 593,...
## $ pctFG           <dbl> 0.519, 0.222, 0.595, 0.357, 0.576, 0.333, 0.376, 0...
## $ fg3m            <dbl> 10, 2, 0, 41, 3, 3, 32, 6, 25, 0, 96, 52, 9, 24, 6,...
## $ fg3a            <dbl> 42, 15, 2, 127, 15, 12, 99, 23, 74, 4, 280, 203, 34...
## $ pctFG3          <dbl> 0.2380952, 0.1333333, 0.0000000, 0.3228346, 0.20000...
## $ pctFT           <dbl> 0.847, 0.700, 0.500, 0.923, 0.735, 0.667, 0.750, 1...
## $ fg2m            <dbl> 674, 2, 481, 15, 277, 10, 35, 5, 13, 3, 161, 669, 1...
## $ fg2a            <dbl> 1277, 3, 807, 30, 471, 27, 79, 13, 36, 6, 313, 1044...
## $ pctFG2          <dbl> 0.5277995, 0.6666667, 0.5960347, 0.5000000, 0.58811...
## $ agePlayer       <dbl> 33, 28, 25, 25, 21, 21, 23, 22, 23, 26, 28, 24, 25,...
## $ minutes         <dbl> 2687, 123, 2669, 588, 1913, 120, 416, 194, 428, 22,...
## $ ftm             <dbl> 349, 7, 146, 12, 166, 8, 45, 4, 7, 1, 150, 500, 37,...
## $ fta             <dbl> 412, 10, 292, 13, 226, 12, 60, 4, 9, 2, 173, 686, 6...
## $ oreb            <dbl> 251, 3, 391, 5, 165, 11, 3, 3, 11, 1, 112, 159, 48,...
## $ dreb            <dbl> 493, 22, 369, 43, 432, 15, 20, 16, 49, 3, 498, 739,...
## $ treb            <dbl> 744, 25, 760, 48, 597, 26, 23, 19, 60, 4, 610, 898,...
## $ ast             <dbl> 194, 8, 124, 20, 184, 13, 25, 5, 65, 6, 104, 424, 1...
## $ stl             <dbl> 43, 1, 117, 17, 71, 1, 6, 1, 14, 2, 68, 92, 54, 22,...
## $ blk             <dbl> 107, 4, 76, 6, 65, 0, 6, 4, 5, 0, 33, 110, 37, 13, ...
## $ tov             <dbl> 144, 4, 135, 14, 121, 8, 33, 6, 28, 2, 72, 268, 58,...
## $ pf             <dbl> 179, 24, 204, 53, 203, 7, 47, 13, 45, 4, 143, 232, ...
## $ pts             <dbl> 1727, 17, 1108, 165, 729, 37, 211, 32, 108, 7, 760,...
## $ urlNBAAPI       <chr> "https://stats.nba.com/stats/playercareerstats?Leag...
## $ n               <int> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, ...
## $ org             <fct> Texas, NA, Other, FC Barcelona Basquet, Kentucky, N...
## $ country         <chr> NA, NA, NA, "Spain", NA, NA, NA, NA, NA, NA, NA, "S...
## $ idConference    <int> 2, 2, 2, 2, 1, 1, 2, 1, 1, 2, 2, 1, 2, 1, 1, 1, ...
```

Continuous

Let's start by taking a look at field goals. It seems pretty likely that this is a continuous variable. Let's take a look at the top 50 spots.

```
nba%>% ## Start with the dataset
  select(namePlayer,slugTeam,fgm)%>% ## and then select a few variables
  arrange(-fgm)%>% ## arrange in reverse order of field goals
  print(n=50) ## print out the top 50
```


```
## # A tibble: 530 × 3
```

##	namePlayer	slugTeam	fgm
##	<chr>	<chr>	<dbl>
##	1 James Harden	HOU	843
##	2 Bradley Beal	WAS	764
##	3 Kemba Walker	CHA	731
##	4 Giannis Antetokounmpo	MIL	721
##	5 Kevin Durant	GSW	721
##	6 Paul George	OKC	707
##	7 Nikola Vucevic	ORL	701
##	8 LaMarcus Aldridge	SAS	684
##	9 Damian Lillard	POR	681
##	10 Karl-Anthony Towns	MIN	681
##	11 Donovan Mitchell	UTA	661
##	12 D'Angelo Russell	BKN	659
##	13 Klay Thompson	GSW	655
##	14 Stephen Curry	GSW	632
##	15 DeMar DeRozan	SAS	631
##	16 Russell Westbrook	OKC	630
##	17 Buddy Hield	SAC	623
##	18 Blake Griffin	DET	619
##	19 Nikola Jokic	DEN	616
##	20 Tobias Harris	MIN	611
##	21 Kyrie Irving	BOS	604
##	22 Devin Booker	PHX	586
##	23 Joel Embiid	PHI	580
##	24 CJ McCollum	POR	571
##	25 Julius Randle	NOP	571
##	26 Andre Drummond	DET	561
##	27 Kawhi Leonard	TOR	560
##	28 LeBron James	LAL	558
##	29 Jrue Holiday	NOP	547
##	30 Montrezl Harrell	LAC	546
##	31 Ben Simmons	PHI	540
##	32 Anthony Davis	NOP	530
##	33 Zach LaVine	CHI	530
##	34 Jordan Clarkson	CLE	529
##	35 Trae Young	ATL	525
##	36 Bojan Bogdanovic	IND	522
##	37 Pascal Siakam	TOR	519
##	38 Collin Sexton	CLE	519
##	39 Jamal Murray	DEN	513
##	40 Deandre Ayton	PHX	509
##	41 Luka Doncic	DAL	506
##	42 Khris Middleton	MIL	506
##	43 De'Aaron Fox	SAC	505
##	44 Andrew Wiggins	MIN	498
##	45 Kyle Kuzma	LAL	496
##	46 Mike Conley	MEM	490
##	47 Lou Williams	LAC	484
##	48 Steven Adams	OKC	481
##	49 Rudy Gobert	UTA	476

```
## 50 Clint Capela          HOU          474  
## # ... with 480 more rows
```


So what I'm seeing here is that field goals aren't "clumped" at certain levels. Let's confirm that by looking at a kernel density plot.

```
nba%>%  
  ggplot(aes(x=fgm)) +  
  geom_density()
```



We can also use a histogram to figure out much the same thing.

```
nba%>%  
  ggplot(aes(x=fgm)) +  
  geom_histogram()
```



Now, technically field goals don't meet the definition I set out above as being a continuous variable because they aren't divisible below a certain amount. Usually in practice though we just ignore this— this variable is “as good as” continuous, given that it varies smoothly over the range and isn't confined to a relatively small set of possible values.

Quick Exercise: Do the same thing for field goal percentage and think about what kind of variable it is.

```
# INSERT CODE HERE
```

Measures for Continuous Variables

The mean is used most of the time for continuous variables, but it's VERY sensitive to outliers. The median (50th percentile) is usually better, but it can be difficult to explain to general audiences.

```
nba%>%
  summarize(mean_fgm=mean(fgm))
```

```
## # A tibble: 1 × 1
##   mean_fgm
##   <dbl>
## 1     191.
```



```
nba%>%
  summarize(median_fgm=median(fgm))
```

```
## # A tibble: 1 × 1
##   median_fgm
##       <dbl>
## 1         157
```

In this case I'd really prefer the mean as a single measure of field goal production, but depending on the audience I still might just go ahead and use the median.

Quick Exercise What measure would you prefer for field goal percentage? Calculate that measure.

```
# INSERT CODE HERE
```

Categorical: ordered

Let's take a look at player seasons.

```
nba%>%
  select(namePlayer,numberPlayerSeason)%>%
  arrange(-numberPlayerSeason)%>%
  print(n=50)
```

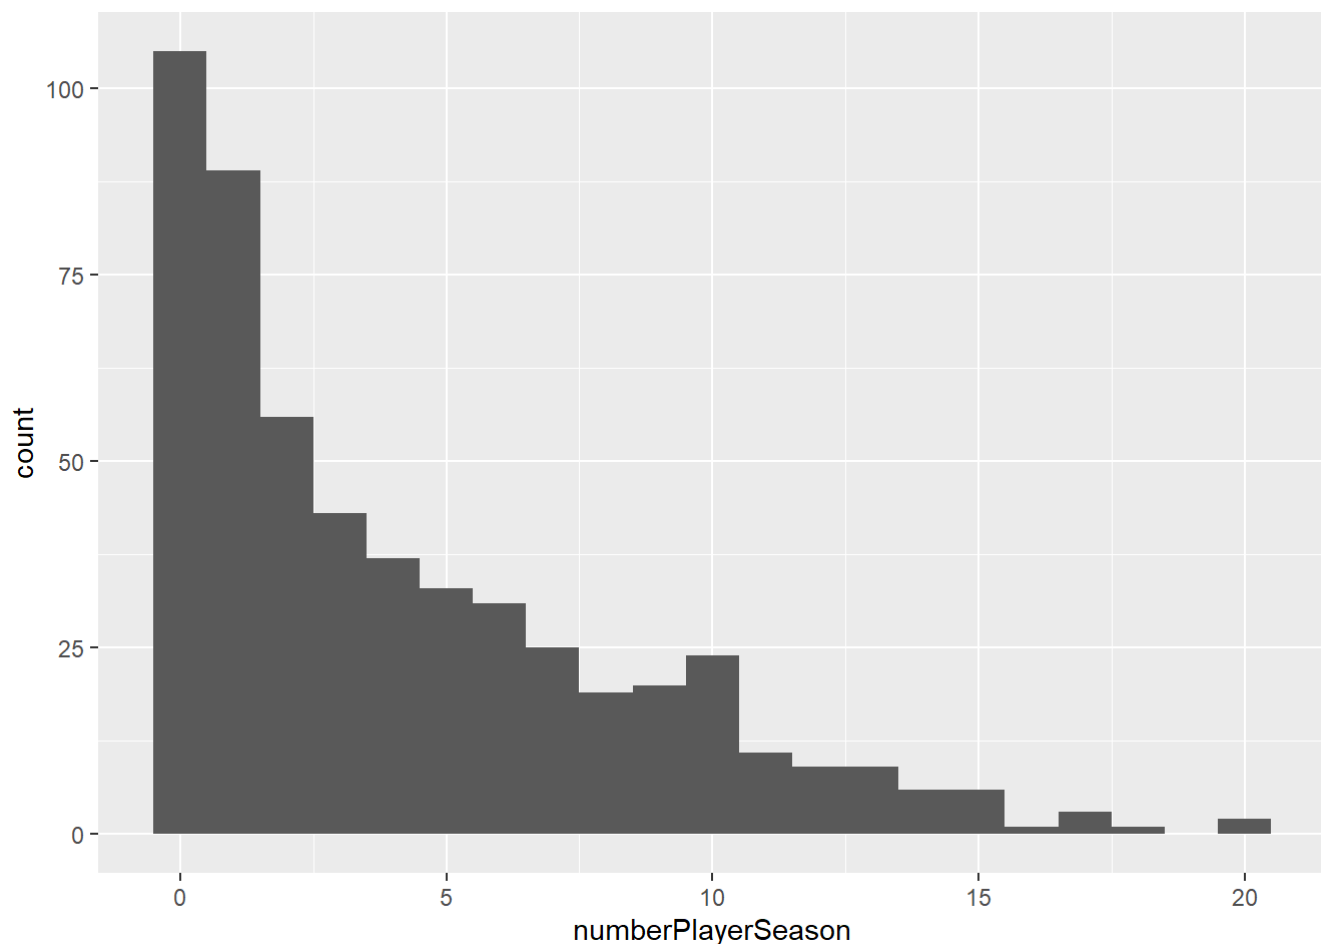
```
## # A tibble: 530 × 2
```

```
##   namePlayer      numberPlayerSeason
##   <chr>                <dbl>
## 1 Vince Carter          20
## 2 Dirk Nowitzki         20
## 3 Jamal Crawford        18
## 4 Tony Parker           17
## 5 Tyson Chandler        17
## 6 Pau Gasol             17
## 7 Nene                  16
## 8 Carmelo Anthony       15
## 9 Udonis Haslem         15
## 10 LeBron James         15
## 11 Zaza Pachulia         15
## 12 Dwyane Wade         15
## 13 Kyle Korver           15
## 14 Luol Deng             14
## 15 Devin Harris          14
## 16 Dwight Howard         14
## 17 Andre Iguodala        14
## 18 JR Smith              14
## 19 Trevor Ariza          14
## 20 Andrew Bogut          13
## 21 Jose Calderon         13
## 22 Raymond Felton        13
## 23 Amir Johnson         13
## 24 Shaun Livingston      13
## 25 Chris Paul            13
## 26 Marvin Williams         13
## 27 Lou Williams          13
## 28 CJ Miles              13
## 29 LaMarcus Aldridge     12
## 30 J.J. Barea            12
## 31 Channing Frye         12
## 32 Rudy Gay              12
## 33 Kyle Lowry            12
## 34 Paul Millsap          12
## 35 JJ Redick             12
## 36 Rajon Rondo           12
## 37 Thabo Sefolosha       12
## 38 Marco Belinelli       11
## 39 Mike Conley           11
## 40 Kevin Durant          11
## 41 Jared Dudley          11
## 42 Marcin Gortat         11
## 43 Gerald Green          11
## 44 Al Horford            11
## 45 Joakim Noah           11
## 46 Thaddeus Young        11
## 47 Nick Young            11
## 48 Corey Brewer          11
## 49 D.J. Augustin         10
```

```
## 50 Jerryd Bayless 10
## # ... with 480 more rows
```

Looks like it might be continuous? Let's plot it:

```
nba%>%
  ggplot(aes(x=numberPlayerSeason))+
  geom_histogram(binwidth = 1)
```



Nope. See how it falls into a small set of possible categories? This is an ordered categorical variable. That means we should calculate the proportions in each category

```
nba%>%
  group_by(numberPlayerSeason)%>%
  count(name="total_in_group")%>%
  ungroup()%>%
  mutate(proportion=total_in_group/sum(total_in_group))
```

```
## # A tibble: 20 × 3
##   numberPlayerSeason total_in_group proportion
##           <dbl>         <int>         <dbl>
## 1             0           105         0.198
## 2             1            89         0.168
## 3             2            56         0.106
## 4             3            43         0.0811
## 5             4            37         0.0698
## 6             5            33         0.0623
## 7             6            31         0.0585
## 8             7            25         0.0472
## 9             8            19         0.0358
## 10            9            20         0.0377
## 11           10            24         0.0453
## 12           11            11         0.0208
## 13           12             9         0.0170
## 14           13             9         0.0170
## 15           14             6         0.0113
## 16           15             6         0.0113
## 17           16             1         0.00189
## 18           17             3         0.00566
## 19           18             1         0.00189
## 20           20             2         0.00377
```

What does this tell us?

Quick Exercise Create a histogram for player age. What does that tell us about the NBA?

```
# INSERT CODE HERE
```

Categorical: ordered, binary

Let's take a look at the variable for Rookie season.

```
nba%>%select(namePlayer,isRookie)
```

```
## # A tibble: 530 × 2
##   namePlayer      isRookie
##   <chr>         <lgl>
## 1 LaMarcus Aldridge FALSE
## 2 Quincy Acy      FALSE
## 3 Steven Adams    FALSE
## 4 Alex Abrines    FALSE
## 5 Bam Adebayo     FALSE
## 6 Rawle Alkins    TRUE
## 7 Grayson Allen   TRUE
## 8 Deng Adel       TRUE
## 9 Jaylen Adams    TRUE
## 10 DeV Vaughn Akoon-Purcell TRUE
## # ... with 520 more rows
```

Okay, so that's set to a logical. In R, TRUE or FALSE are special values that indicate the result of a logical question. In this it's whether or not the player is a rookie.

Usually we want a binary variable to have at least one version that's structured so that 1= TRUE and 2=FALSE. This makes data analysis much easier. Let's do that with this variable.

This code uses `ifelse` to create a new variable called `isRookiebin` that's set to 1 if the `isRookie` variable is true, and 0 otherwise.

```
nba<-nba%>%
  mutate(isRookie_bin=ifelse(isRookie==TRUE,1,0))
```

Now that it's coded 0,1 we can calculate the mean, which is the same thing as the proportion of the players that are rookies.

```
nba%>%summarize(mean=mean(isRookie_bin))
```

```
## # A tibble: 1 × 1
##   mean
##   <dbl>
## 1 0.198
```

Categorical: unordered

Let's take a look at which college a player attended, which is a good example of an unordered categorical variable. The `org` variable tells us which organization the player was in before playing in the NBA.

```
nba%>%
  select(org)%>%
  glimpse()
```

```
## Rows: 530
## Columns: 1
## $ org <fct> Texas, NA, Other, FC Barcelona Basquet, Kentucky, NA, Duke, NA, NA...
```

This look like team or college names, so this would be a categorical variable. Let's take a look at the counts of players from different organizations:

```
nba%>%
  group_by(org)%>%
  count()%>%
  arrange(-n)%>%
  print(n=50)
```

```
## # A tibble: 68 × 2
## # Groups:   org [68]
##   org          n
##   <fct>      <int>
## 1 <NA>        157
## 2 Other         85
## 3 Kentucky      25
## 4 Duke          17
## 5 California-Los Angeles 15
## 6 Kansas         11
## 7 Arizona        10
## 8 Texas          10
## 9 North Carolina   9
## 10 Michigan        8
## 11 Villanova       7
## 12 Indiana         6
## 13 Southern California  6
## 14 Syracuse         6
## 15 California       5
## 16 Louisville       5
## 17 Ohio State       5
## 18 Wake Forest      5
## 19 Colorado         4
## 20 Connecticut      4
## 21 Creighton        4
## 22 FC Barcelona Basquet 4
## 23 Florida          4
## 24 Georgia Tech     4
## 25 Michigan State   4
## 26 Oregon           4
## 27 Utah             4
## 28 Washington       4
## 29 Wisconsin        4
## 30 Boston College   3
## 31 Florida State    3
## 32 Georgetown       3
## 33 Gonzaga          3
## 34 Iowa State       3
## 35 Marquette        3
## 36 Maryland         3
## 37 Miami (FL)       3
## 38 North Carolina State 3
## 39 Notre Dame       3
## 40 Oklahoma         3
## 41 Purdue           3
## 42 Southern Methodist 3
## 43 Stanford         3
## 44 Tennessee        3
## 45 Virginia         3
## 46 Anadolu Efes S.K.  2
## 47 Baylor           2
## 48 Butler           2
```

```
## 49 Cincinnati                2
## 50 Kansas State              2
## # ... with 18 more rows
```

Here we have a problem. If we're interested just in colleges, we're going to need to structure this a bit more. The code below filters out three categories that we don't want: missing data, anything classified as others, and sports teams from other countries. The last is incomplete— I probably missed some! If I were doing this for real, I would use a list of colleges and only include those names.

What I do below is to negate the `str_detect` variable by placing the `!` in front of it. This means I want all of the cases that don't match the pattern the supplied. The pattern makes heavy use of the OR operator `|`. I'm saying I don't want to include players whose organization included the letters `CB r KK` and so on (these are common prefixes for sports organizations in other countries, I definitely did not look that up on Wikipedia. Ok, I did.).

```
nba%>%
  filter(!is.na(org))%>%
  filter(!org=="Other")%>%
  filter(!str_detect(org, "CB|KK|rytas|FC|B.C.|S.K.|Madrid"))%>%
  group_by(org)%>%
  count()%>%
  arrange(-n)%>%
  print(n=50)
```

```
## # A tibble: 57 × 2
## # Groups:   org [57]
##   org          n
##   <fct>      <int>
## 1 Kentucky      25
## 2 Duke           17
## 3 California-Los Angeles 15
## 4 Kansas         11
## 5 Arizona        10
## 6 Texas          10
## 7 North Carolina   9
## 8 Michigan         8
## 9 Villanova        7
## 10 Indiana         6
## 11 Southern California  6
## 12 Syracuse        6
## 13 California       5
## 14 Louisville       5
## 15 Ohio State       5
## 16 Wake Forest     5
## 17 Colorado        4
## 18 Connecticut     4
## 19 Creighton       4
## 20 Florida         4
## 21 Georgia Tech    4
## 22 Michigan State  4
## 23 Oregon          4
## 24 Utah            4
## 25 Washington      4
## 26 Wisconsin       4
## 27 Boston College  3
## 28 Florida State   3
## 29 Georgetown      3
## 30 Gonzaga         3
## 31 Iowa State      3
## 32 Marquette       3
## 33 Maryland        3
## 34 Miami (FL)      3
## 35 North Carolina State 3
## 36 Notre Dame      3
## 37 Oklahoma        3
## 38 Purdue          3
## 39 Southern Methodist 3
## 40 Stanford        3
## 41 Tennessee       3
## 42 Virginia        3
## 43 Baylor          2
## 44 Butler          2
## 45 Cincinnati      2
## 46 Kansas State    2
## 47 Louisiana State 2
## 48 Memphis         2
```



```
## 49 Missouri                2
## 50 Murray State            2
## # ... with 7 more rows
```

That looks better. Which are the most common colleges and universities that send players to the NBA?

Quick Exercise Arrange the number of players by team in descending order.

```
# INSERT CODE HERE
```

Categorical: unordered, binary

There are two conference in the NBA, eastern and western. Let's take a look at the variable that indicates which conference the player played in that season.

```
nba%>%select(idConference)%>%
  glimpse()
```

```
## Rows: 530
## Columns: 1
## $ idConference <int> 2, 2, 2, 2, 1, 1, 2, 1, 1, 2, 2, 1, 2, 1, 1, 1, 1, 2, 2, ...
```

It looks like conference is structured as numeric, but a “1” or a “2”. Because it's best to have binary variables structured as “has the characteristic” or “doesn't have the characteristic” we're going to create a variable for western conference that's set to 1 if the player was playing in the western conference and 0 if the player was not (this is the same as playing in the eastern conference).

```
nba<-nba%>%
  mutate(west_conference=ifelse(idConference==1,1,0))
```

Once we've done that, we can see how many players played in each conference.

```
nba%>%
  summarize(mean(west_conference))
```

```
## # A tibble: 1 × 1
##   `mean(west_conference)`
##               <dbl>
## 1                   0.508
```

Makes sense!

Quick Exercise: create a variable for whether or not the player is from the USA. Calculate the proportion of players from the USA in the NBA. The coding on country is ... decidedly US-centric, so you'll need to think about this one a bit.

```
# INSERT CODE HERE
```

Analysis

Ok, now that we know how this works, we can do some summary analysis. First of all, what does the total number of field goals made look like by college?

We know that field goals are continuous (sort of) so let's summarize them via the mean. We know that college is a categorical variable, so we'll use that to group the data. This is one of our first examples of a conditional mean, which we'll use a lot.

Top 50 Colleges by Total FG

```
nba%>%
  filter(!is.na(org))%>%
  filter(!org=="Other")%>%
  filter(!str_detect(org, "CB|KK|rytas|FC|B.C.|S.K.|Madrid"))%>%
  group_by(org)%>%
  summarize(mean_fg=sum(fgm))%>%
  arrange(-mean_fg)%>%
  print(n=50)
```

```
## # A tibble: 57 × 2
##   org                                mean_fg
##   <fct>                             <dbl>
## 1 Kentucky                          6594
## 2 Duke                              4623
## 3 Texas                             3437
## 4 California-Los Angeles            3382
## 5 Kansas                            2765
## 6 Arizona                           2101
## 7 Oklahoma                          1767
## 8 Southern California               1758
## 9 Louisville                         1679
## 10 North Carolina                   1659
## 11 Indiana                          1522
## 12 Ohio State                       1486
## 13 Michigan                         1392
## 14 Wake Forest                      1364
## 15 Connecticut                      1299
## 16 Villanova                        1222
## 17 Georgia Tech                     1169
## 18 Tennessee                        1095
## 19 Stanford                          949
## 20 Utah                             943
## 21 Marquette                        873
## 22 Gonzaga                          863
## 23 Michigan State                   820
## 24 Colorado                         818
## 25 Virginia                         816
## 26 Maryland                         811
## 27 Missouri                         756
## 28 California                       734
## 29 Florida State                    733
## 30 Georgetown                       717
## 31 Memphis                          620
## 32 Florida                          618
## 33 North Carolina State             598
## 34 Boston College                   586
## 35 Louisiana State                  583
## 36 Syracuse                         567
## 37 Iowa State                       523
## 38 Butler                           459
## 39 Wisconsin                        456
## 40 Creighton                       432
## 41 Oregon                           352
## 42 Texas A&M                        322
## 43 Baylor                           312
## 44 Providence                       291
## 45 Purdue                           275
## 46 Notre Dame                       263
## 47 Ulkerspor                        252
## 48 Southern Methodist               246
## 49 Oklahoma State                   242
```

```
## 50 West Virginia                236
## # ... with 7 more rows
```

Next, what about field goal percentage?

Top 50 Colleges by Average Field Goal Percent

```
nba%>%
  filter(!is.na(org))%>%
  filter(!org=="Other")%>%
  filter(!str_detect(org, "CB|KK|rytas|FC|B.C.|S.K.|Madrid"))%>%
  group_by(org)%>%
  summarize(mean_ftp=mean(pctFT))%>%
  arrange(-mean_ftp)%>%
  print(n=50)
```

```
## # A tibble: 57 × 2
##   org                mean_ftp
##   <fct>              <dbl>
## 1 Tennessee          0.842
## 2 Virginia            0.833
## 3 Oklahoma            0.823
## 4 North Carolina State 0.817
## 5 West Virginia       0.804
## 6 Ulkerspor           0.803
## 7 Missouri            0.802
## 8 Wake Forest         0.802
## 9 Florida State       0.801
## 10 Murray State       0.798
## 11 Iowa State         0.795
## 12 Notre Dame         0.792
## 13 Memphis            0.788
## 14 Florida             0.784
## 15 Michigan            0.783
## 16 Stanford            0.779
## 17 Georgetown          0.775
## 18 Marquette           0.774
## 19 Utah                0.770
## 20 Kansas State       0.767
## 21 Butler              0.762
## 22 Gonzaga             0.761
## 23 North Carolina     0.756
## 24 Villanova           0.755
## 25 Texas               0.752
## 26 Connecticut        0.748
## 27 Providence          0.747
## 28 Boston College     0.742
## 29 Michigan State     0.730
## 30 Kansas              0.729
## 31 Indiana             0.729
## 32 Duke                0.728
## 33 Baylor              0.726
## 34 Arizona             0.721
## 35 Pallacanestro Biella 0.718
## 36 Wisconsin          0.712
## 37 Kentucky            0.712
## 38 Georgia Tech        0.712
## 39 Louisiana State    0.709
## 40 Creighton          0.698
## 41 Maryland            0.695
## 42 Vanderbilt          0.688
## 43 Washington          0.680
## 44 Louisville          0.679
## 45 Ohio State         0.679
## 46 California          0.675
## 47 Southern Methodist 0.673
## 48 Oregon              0.662
## 49 Texas A&M           0.652
```

```
## 50 Southern California      0.648
## # ... with 7 more rows
```

Quick Exercise Calculate field goals made by player season.

```
# INSERT CODE HERE
```

Quick Exercise Calculate free throw percent made by player season.

```
# INSERT CODE HERE
```